

Operating Systems

سیستم عامل

مهدی کارگهی
ترم اول ۹۷-۱۳۹۶

References

■ **A. Silberschatz, P. B. Galvin, and G. Gagne, *Operating System Concepts*, 9th Ed., 2013. (10th Ed., Feb. 2018)**

■ Key: OS96971MK

■ Email ():

■ Start: Before 9:10AM

■ Simultaneous: Course+Lab (4 Units)

■ Cutoff

■ Midterm: 30% (Date:)

■ Final: 30%

■ Projects: 7-8 prjs. 30%

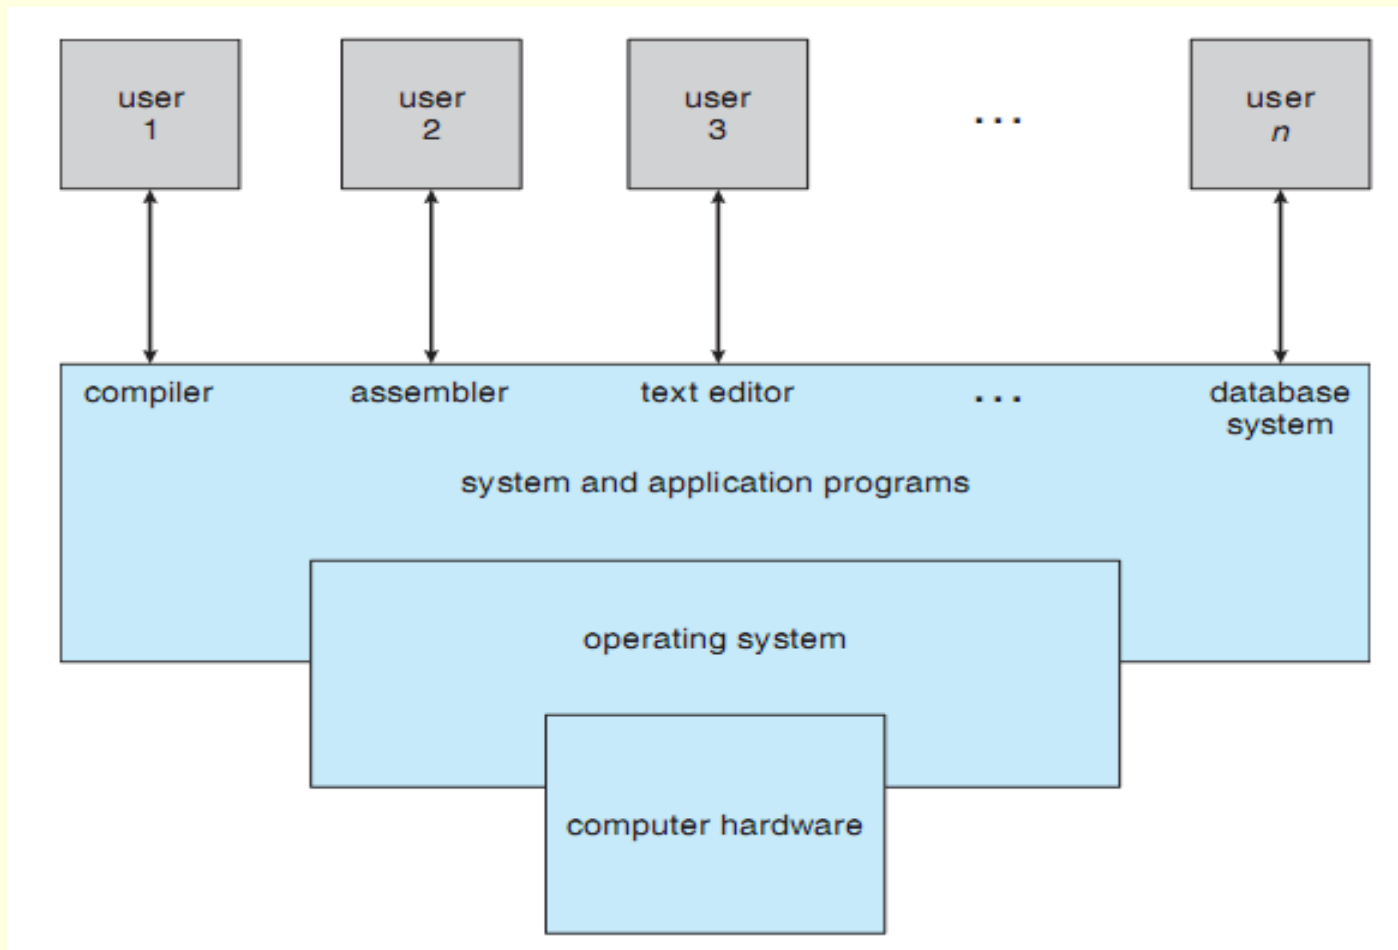
■ Quizzes & Assignments: Up to 10%

■ TA classes (Date & Time:)

Course Outline

- Introduction (Terminologies & Overview)
- Different Structures of Operating Systems
- Processes and Threads
- Process Synchronization: Facts and Mechanisms
- CPU Scheduling (Approaches and Algorithms)
- Resource Sharing and Deadlock
- Main and Virtual Memory Management
- Storage Management
- Protection, Security, and Special Purpose Systems

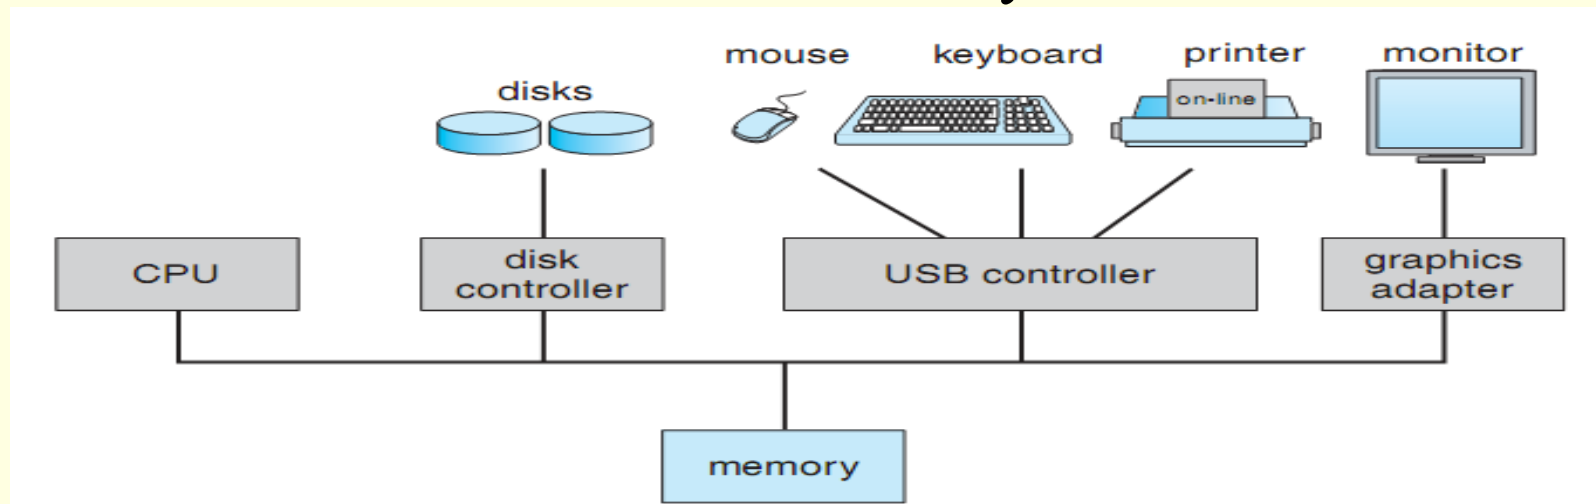
Components of a Computer System





Computer-System Organization


- Firmware: Bootstrap program (in ROM or EEPROM)
 - Initializes CPU registers, device controllers, memory contents
 - Locates and loads into memory the OS kernel



How the OS Discovers the Occurrence of an Event

- Polling
- Interrupt
 - From hardware by a signal sent to the CPU
 - From software by system call (monitor call)

Interrupt time line for a single process doing output



Polling vs. Interrupt

■ **Interrupts are used when:**

- Efficiency is paramount (time is an important resource), or
- Multiple devices must be monitored simultaneously (more concurrency).

■ **Polling is typically used when:**

- The processor must respond to some event more quickly than is possible using interrupts, or
- Large amounts of data are expected to arrive at particular intervals, such as during real-time data acquisition

How an Interrupt Triggers?

1. Hardware Interrupt

- Some external HW changes the voltage level on an int. req. line
- Once the ISR completes, the program resumes where it left off

2. Software Interrupt

- The program that is executing triggers the interrupt by
 - executing a special instruction or
 - writing to a memory-mapped register
- Once the ISR completes, the program resumes where it left off

3. Exception

- The interrupt is triggered by internal HW that detects a fault, such as a segmentation fault
- Once the ISR has completed, the program is **not** normally resumed. Instead, the PC is set to some fixed location where, e.g., the OS may terminate the offending program.
- Typically, exceptions have the highest priority, serviced always.

Interrupts

Program memory addresses,
not data memory addresses.

The most typical and general program setup for the Reset and Interrupt Vector Addresses in ATmega168 is:

Address	Labels	Code	Comments
0x0000		jmp RESET	; Reset Handler
0x0002		jmp EXT_INT0	; IRQ0 Handler
0x0004		jmp EXT_INT1	; IRQ1 Handler
0x0006		jmp PCINT0	; PCINT0 Handler
0x0008		jmp PCINT1	; PCINT1 Handler
0x000A		jmp PCINT2	; PCINT2 Handler
0x000C		jmp WDT	; Watchdog Timer Handler
0x000E		jmp TIM2_COMPA	; Timer2 Compare A Handler
0x0010		jmp TIM2_COMPB	; Timer2 Compare B Handler
0x0012		jmp TIM2_OVF	; Timer2 Overflow Handler
0x0014		jmp TIM1_CAPT	; Timer1 Capture Handler

Source: ATmega168 Reference Manual

Responses:

- Disable interrupts.
- Push the current PC onto the stack.
- Execute the instruction at a designated address in the flash memory.

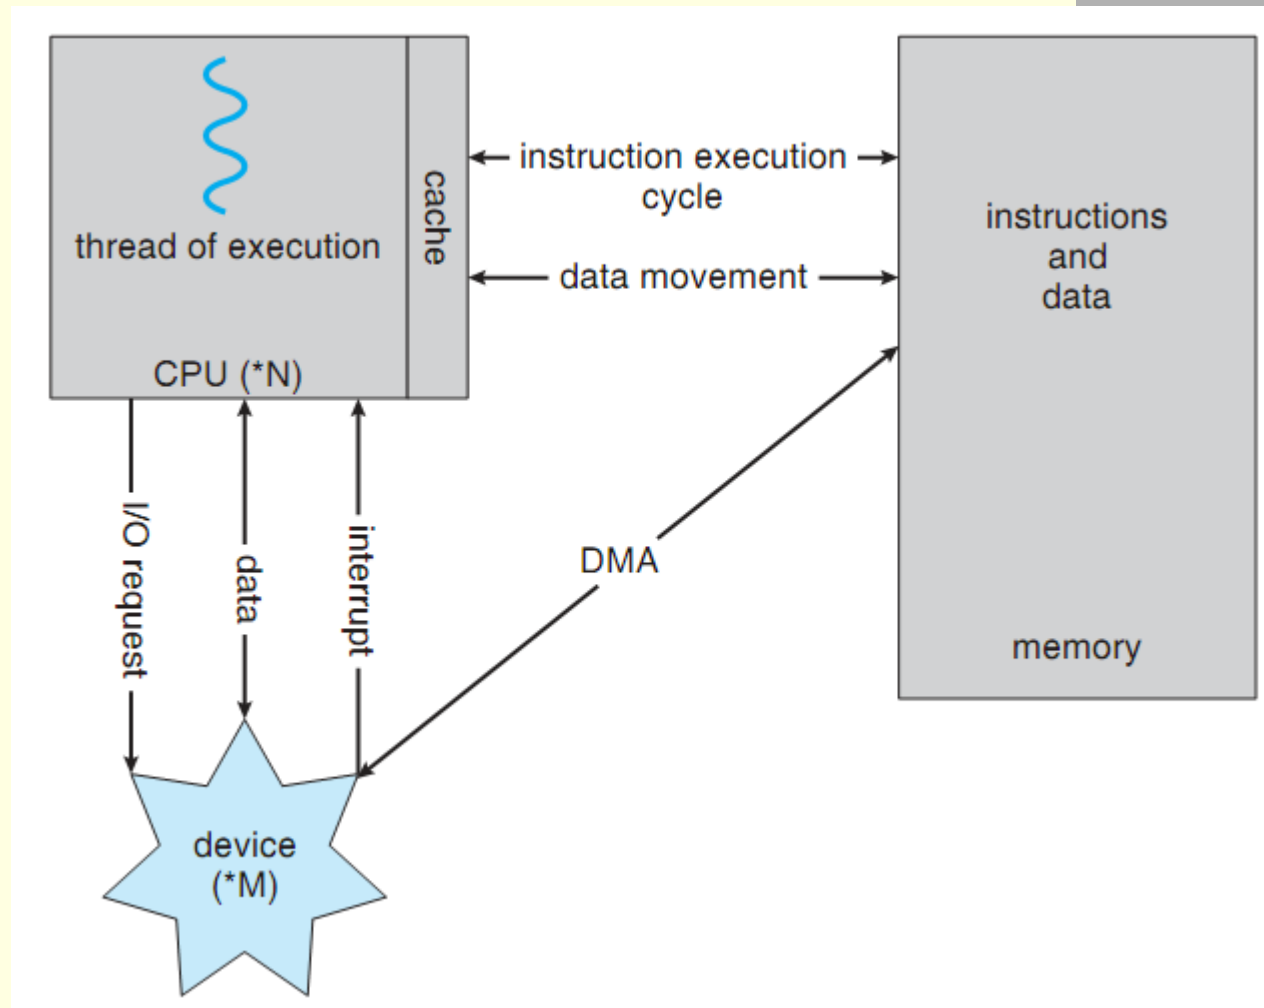
Design of ISR:

- Save and restore any registers it uses.
- Re-enable interrupts before returning from interrupt.

Interrupts

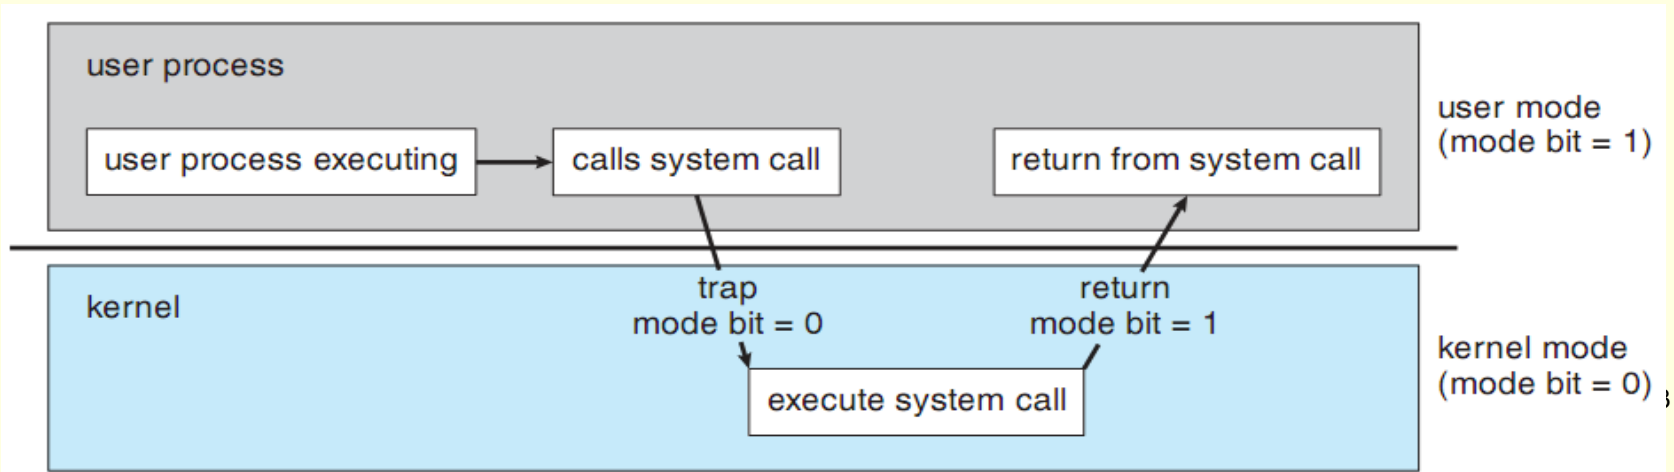
- Interrupts inform the completion of I/O to the CPU
- Interrupt Service Routine (ISR)
- Device drivers?
- Typically: A device driver for each device controller!
- Examples
- Direct Memory Access (DMA)
- What OS does with respect to DMA?
 - Determination of buffers, pointers, and counters for the I/O device
 - Sending the start command
 - The Completion of transfer will be notified by an interrupt
- DMA steals memory cycles from CPU to access the bus

The Interplay of Components of a Computer System



Operating-System Operations

- Modern operating systems are *interrupt-driven*
- *Trap* is a software interrupt
- Dual-Mode operation (through mode bit)
 - User mode (1)
 - Kernel (Supervisor, System, or Privileged) mode (0)
 - Privileged instructions



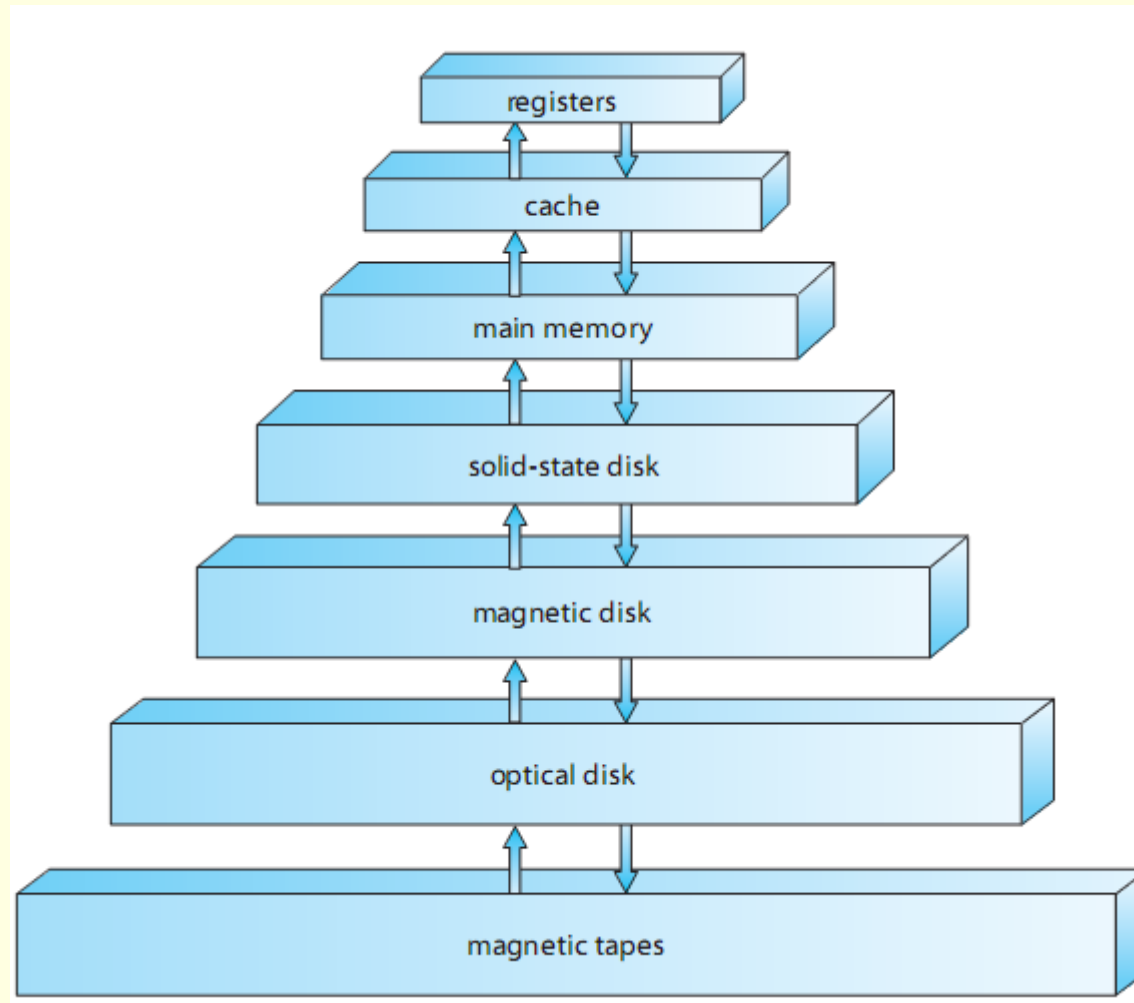
Privileged Instructions

- The instruction to change mode bit
- I/O instructions
- Timer management
- Interrupt (and interrupt vector) management
 - What would occur if it was not privileged?
- ...
- Errors violating modes are detected by the **hardware** and it will trap to the OS. The trap transfers control through the interrupt vector to the OS.

Timer

- Fixed or variable timer
- Preventing a program from getting stuck in an infinite loop
- Time sharing

Storage Hierarchy



Multi-Processor (Tightly-Coupled) Systems

- Sharing



- the computer bus, sometimes the clock, memory, and peripheral devices

- Main advantages

- Increased throughput: the speedup with N processors is not N due to overheads, contention for shared resources, and that the degree of requests is not N times of a single-processor system
 - Economy of scale: due to sharing of resources
 - Increased reliability: just slows down due to failure
 - Graceful degradation
 - Fault tolerance: differs from graceful degradation

Multi-Processor (Tightly-Coupled) Systems

- Two types (HW/SW)
 - Asymmetric multiprocessing
 - Master-slave relationship (e.g., SunOS Ver. 4)
 - Symmetric multiprocessing (SMP)
 - E.g., Windows, Linux, and Mac OS
 - Multi-core CPUs



Blade servers

- Multiple processor boards, I/O boards, and networking boards on the same chassis
- They boot independently and run their own OS



Operating-System Structure

- Program vs. process/task/job
- Batch systems
- Buffering
- Spooling
- Multiprogramming
 - Increasing CPU utilization
 - Jobs are loaded into memory → Scheduling & Memory management
 - Switching at *job completion/wait for IO*
- Time-sharing (Multitasking)
 - User interaction with the system
 - Switching between users and tasks with high freq.
 - Switching at *job completion/wait for IO/completion of time-slice*

Operating-System Structure

- Job scheduling
 - Selecting a job from the job pool to be loaded into memory
- CPU scheduling
 - Selecting the next ready job to run
- Virtual memory
 - Tries to separate logical memory from physical memory through swapping
- Time sharing may require CPU sch., mem. and virtual mem. managements, file system, disk management, protection of resources, job synchronization, deadlock management

Operating-System Operations

- Process Management
- Memory Management
- Storage Management
 - File-System
 - Mass-Storage
 - Caching
 - I/O Systems
- Protection and Security
- Command Interpreter

Process Management

- What is a process?
- Processes require resources
- Processes are created by other processes
- What about the first process?
 - Creating and deleting both user and system processes
 - Suspending and resuming processes
 - Providing mechanisms for process synchronization
 - Providing mechanisms for process communication
 - Providing mechanisms for deadlock handling

Memory Management

- Memory is a large array of words
- Keeping track of which parts of memory are currently being used and by whom
- Deciding which processes and data to move into and out of memory
- Allocating and deallocating memory space as needed

Storage Management

- The capacity of main memory is restricted →
Some kinds of support is required
- Mass-Storage Management
 - Free-space management
 - Storage allocation
 - Disk scheduling
- Secondary storage
- Tertiary storage (slower and cheaper)
 - WORM (Write-Once Read-Many)
 - RW (Read-Write)

Storage Management

- File is a logical storage unit
- File-System Management
 - Creating and deleting files
 - Creating and deleting directories to organize files
 - Supporting primitives for manipulating files and directories
 - Mapping files onto secondary storage
 - Backing up files on stable (nonvolatile) storage media

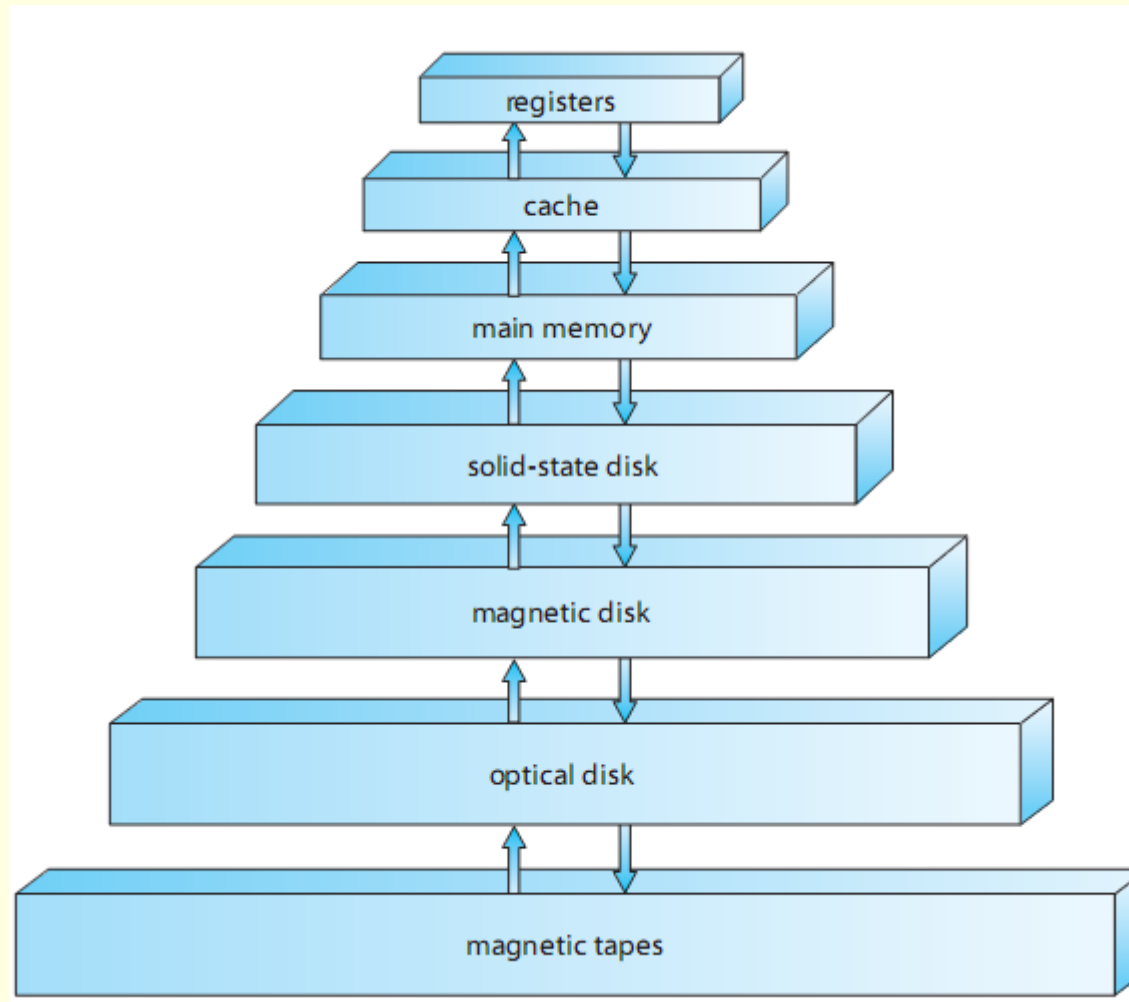
Storage Management

- I/O Systems
 - The purpose of OS is to hide the details of I/O operations from the user
- I/O subsystem
 - A memory management component that includes buffering, caching, and spooling
 - A general device-driver interface
 - Drivers for specific hardware devices
- Only the device driver knows the details of the behavior of the specific device

Protection and Security

- Protection: Any mechanism for controlling the access of processes or users to the resources
 - Distinguishing authorized and unauthorized users
- Security: Defending a system from external and internal attacks (e.g., viruses, worms, denial of service attacks, etc).
 - A system may be protected, but not secure
 - It should be done by the OS and additional software and/or policies

Storage Hierarchy



Performance of Various Levels of Storage

Level	1	2	3	4	5
Name	registers	cache	main memory	solid state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25 - 0.5	0.5 - 25	80 - 250	25,000 - 50,000	5,000,000
Bandwidth (MB/sec)	20,000 - 100,000	5,000 - 10,000	1,000 - 5,000	500	20 - 150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape

Storage Management

- Cache Management
 - Careful selection of cache size
 - Replacement policy
- Different levels of cache
 - Controlled by OS (disk to memory) or HW (memory to registers)
- Multiprocessor systems
 - Cache coherency
- Distributed systems
 - Replicated files

Computing Environments

- Traditional (Time Sharing) Computing: Inline
- Mobile Computing
- Distributed Systems
- Client/Server Computing
- Peer-to-Peer Computing
- Virtualization
- Cloud Computing
- Real-Time Embedded Systems

Mobile Computing


- Such as PDAs, cellular phones, ...
- Many of them use specific operating systems: iOS and Android
- Properties:
 - small memory, (512KB to a few GB)→require efficient memory management
 - Slow processors (power & size limitations)
 - Small keyboard and display (I/O)→ web clipping
 - Wireless technology (BlueTooth or 802.11)
 - GPS, Gyroscope

Distributed Systems

- A network (LAN/MAN/WAN/etc) of physically separate, possibly heterogeneous computer systems
 - Computation speedup
 - Resource sharing
 - Increasing functionality and data availability
 - Increasing reliability
- Network OS (NOS) vs. Distributed OS (DOS)
- NFS/DSM/...


Client-server computing

- Servers are bottlenecks
 - Compute-server systems (e.g., database servers)
 - File-server systems (e.g., web servers)



Peer-to-peer computing

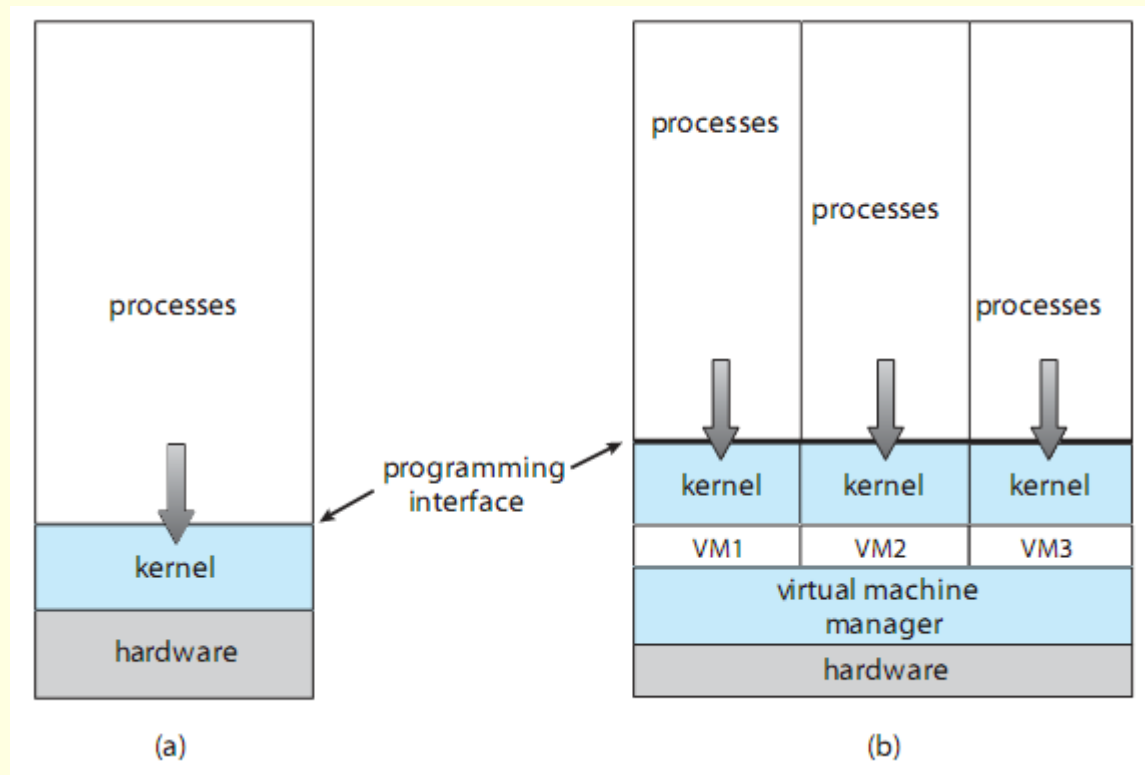
- Each peer may be a client or a server
 - Idea-1: A node joins to network, registers its service, any other node asks a centralized lookup service, two peers communicate each other as a C/S system (e.g., Napster file-sharing service)
 - Idea-2: The client peer Bcasts a request for a desired service, the nodes providing the service respond to the peer, ... (e.g., Gnutella file-sharing service)



Virtualization

- Allows operating systems to run applications within other OSes
 - Vast and growing industry
- **Emulation** used when source CPU type different from target type (i.e. PowerPC to Intel x86)
 - Generally slowest method
 - When computer language not compiled to native code – **Interpretation**
- **Virtualization** – OS natively compiled for CPU, running **guest** OSes also natively compiled
 - Consider VMware running WinXP guests, each running applications, all on native WinXP **host** OS
 - **VMM** provides virtualization services

Virtualization



Cloud Computing

- Delivers computing, storage, even apps as a service across a network
- Logical extension of virtualization as based on virtualization
 - Amazon Elastic Compute Cloud (**EC2**) has thousands of servers, millions of VMs, petabytes of storage available across the Internet, pay based on usage
- Many types
 - **Public cloud** – available via Internet to anyone willing to pay
 - **Private cloud** – run by a company for the company's own use
 - **Hybrid cloud** – includes both public and private cloud components
 - Software as a Service (**SaaS**) – one or more applications available via the Internet (i.e. word processor)
 - Platform as a Service (**PaaS**) – software stack ready for application use via the Internet (i.e. a database server)
 - Infrastructure as a Service (**IaaS**) – servers or storage available over Internet (i.e. storage available for backup use)

Real-Time Embedded Systems

- Embedded systems
 - Special purpose
 - Little or no UI
- Examples: car engines, manufacturing robots, VCRs, heating and lighting, ...
- They may use general-purpose OS (e.g., UNIX), special-purpose embedded OS (e.g., Windows CE), or application-specific integrated circuits (ASICs) with no OS
- Their use continues to expand rapidly!
- Power consumption is a main concern!

Real-Time Embedded Systems

- Almost always are time-sensitive in the process of sensing/computing/actuating
- Examples: medical imaging, industrial control, automobile-engine fuel-injection, home appliances, military systems
- Real-time systems have two properties
 - Computationally correctness
 - Temporally correctness
- HRT vs. SRT
- Therefore, they need real-time operating systems (RTOS)
- Specialized scheduling/memory management/...
- Secondary storage?