

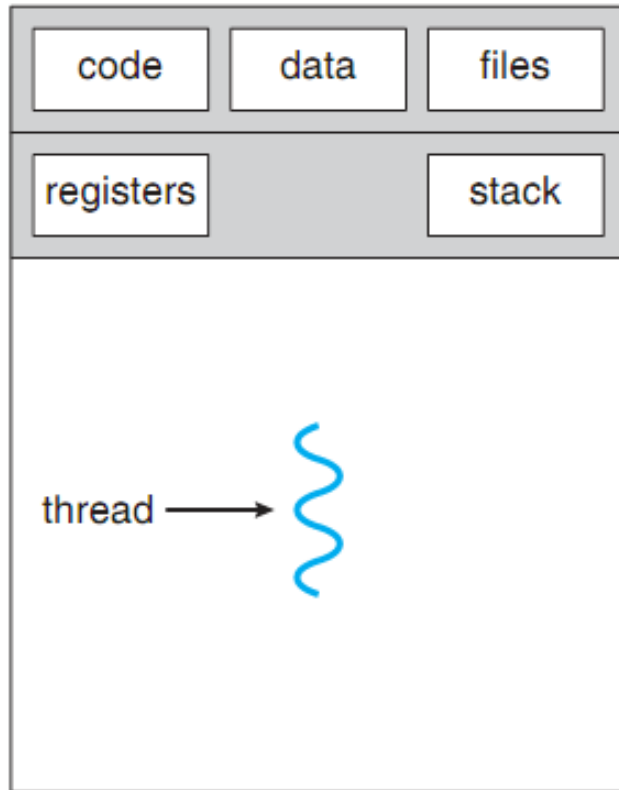
Threads

Mehdi Kargahi
School of ECE
University of Tehran
Summer 2016

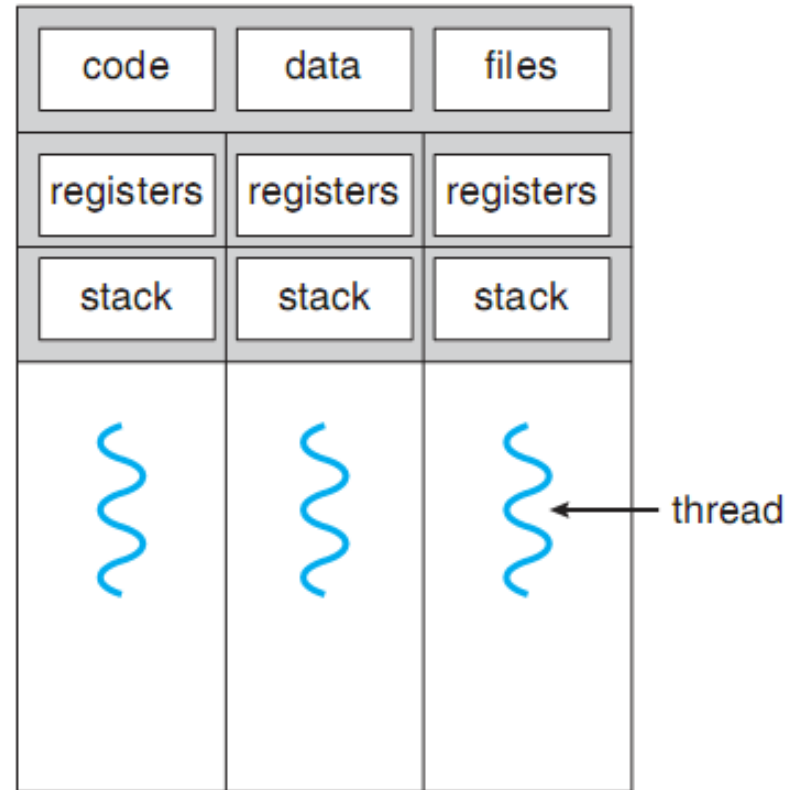
Definition

- What's the difference between a thread and a process?
- Heavyweight process: A single thread process
- What a thread has for itself?
 - Thread-Specific Data
 - Thread ID
 - Program Counter
 - A register set
 - A stack

Single-threaded and Multithreaded Processes



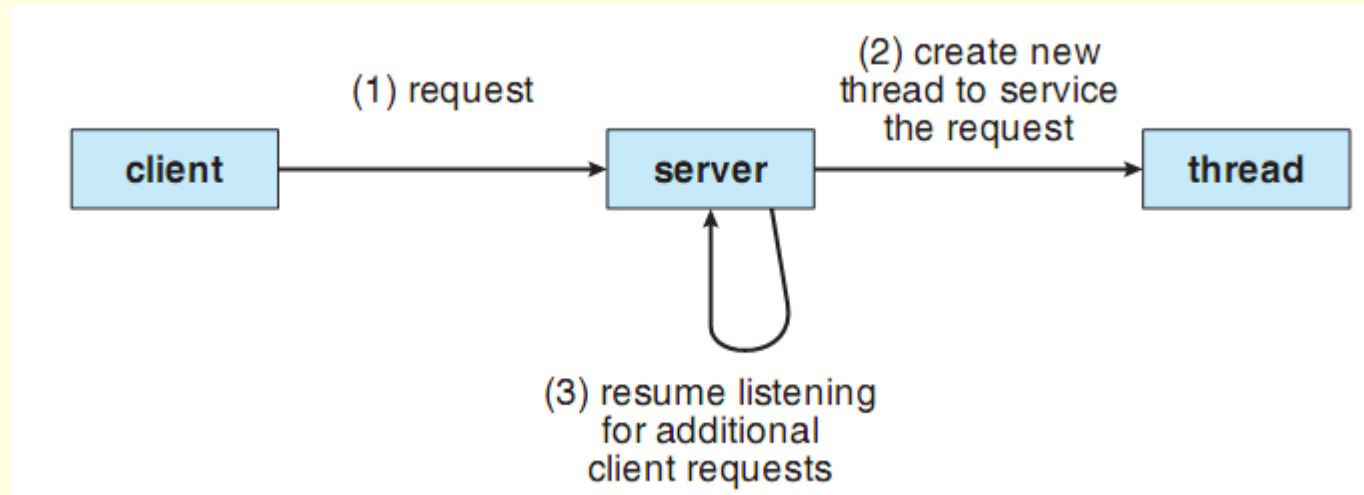
single-threaded process



multithreaded process

Examples

- A Web server
- A RPC server
- Since the requests from these servers are similar, it is better to avoid creating new processes that may be much more costly



The Benefits of Multithreaded Programming

- Responsiveness
 - If a thread is blocked, other threads can continue
 - E.g., showing an image while getting text from the user in a web page
- Resource sharing
 - Threads share the memory and resources of the respective process (shared code and data)
- Economy
 - Allocating memory and resource to processes and their respective context-switch is more costly with respect to threads
 - Solaris: speed of process creation=1/30 speed of thread creation
 - Solaris: speed of process CS=1/5 speed of thread CS
- Utilization of multiprocessor architectures
 - No speed-up with a heavyweight process on a multiprocessor system
 - Concurrency with multithreaded processes on multiprocessor systems

Multicore Programming

- Amdahl's Law

- S: The portion of the application that must be performed serially
- N: the number of processing cores
 - Suppose that N goes to infinity!

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

Challenges in Programming for Multicore Systems

- **Identifying tasks:** Finding areas that can be divided into separate, concurrent tasks.
- **Balance:** Programmers must ensure that the tasks perform equal work of equal value to assign a core for which.
- **Data splitting:** The data accessed and manipulated by the tasks must also be divided to run on separate cores.
- **Data dependency:** When one task depends on data from another, programmers must ensure that the execution of the tasks is synchronized to accommodate the data dependency.
- **Testing and debugging:** When a program is running in parallel on multiple cores, many different execution paths are possible. Testing and debugging such concurrent programs is inherently more difficult.

Types of Parallelism

- **Data Parallelism:** Distributing subsets of the same data across multiple computing cores and performing the same operation on each core.
 - Example: Summing the contents of an array of size N .
- **Task Parallelism:** Distributing not data but tasks (threads) across multiple computing cores. Each thread is performing a unique operation.

In most instances, applications use a hybrid of these two strategies.

Multithreading Models

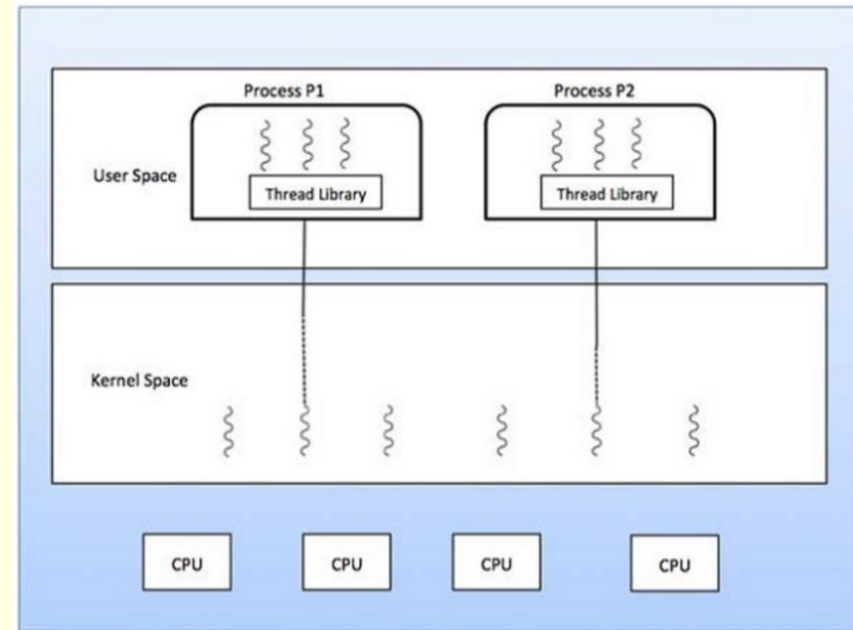
- *User threads* are supported above the kernel
- *Kernel threads* are supported and managed by the OS

Thread Libraries

- Pthreads ([Book::P.173](#))
- Win32 threads
- Java threads

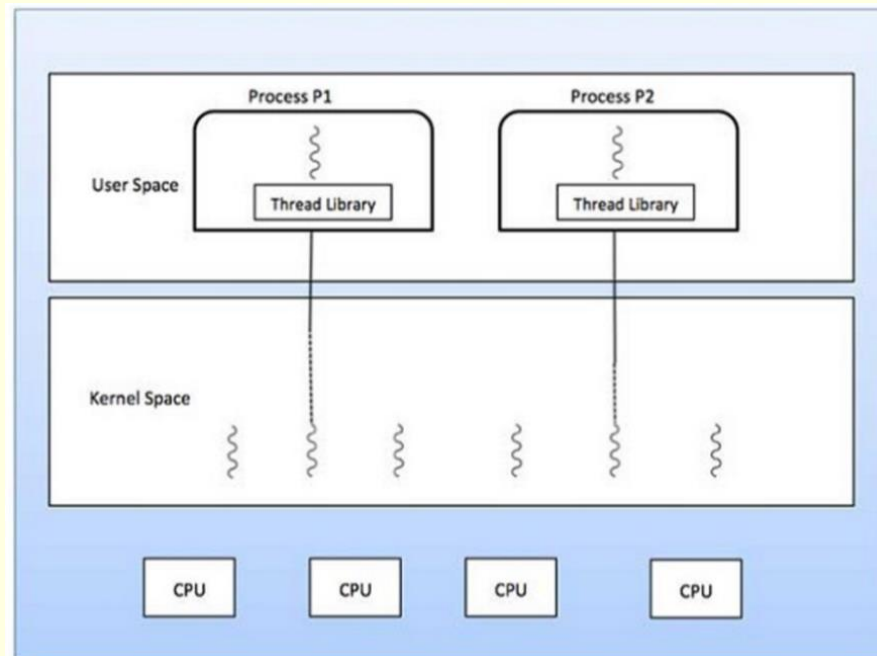
Many-to-One Model

- Benefit
 - Since thread management is in user space, it is efficient
- Disadvantage
 - A blocking system call will block the entire process
 - No benefit of using a multiprocessor system
- Examples
 - Green threads in Solaris
 - GNU Portable threads



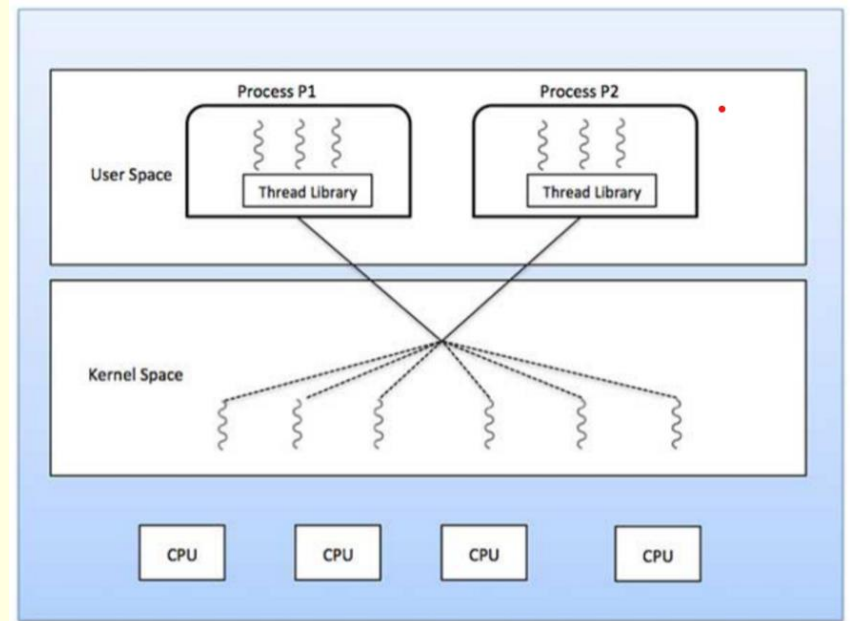
One-to-One Model

- Benefits
 - A blocking system call will not block the entire process
 - Multiple threads can run concurrently on a MP system
- Disadvantage
 - Creating a user thread requires creating a corresponding kernel thread which is costly and usually restricts the number of threads supported by the system
- Examples
 - Linux and the family of Windows




Many-to-Many Model

- Multiplexing (Tru64 UNIX)
 - No. of kernel threads \leq no. of user threads
- Properties
 - As many user threads as necessary can be created
 - The corresponding kernel threads can run in parallel on a multiprocessor
 - When a thread performs a blocking system call, the kernel can schedule another thread



Two-Level Model



Differences Between User-Level and Kernel-Level Threads

S.N.	User-Level Threads	Kernel-Level Thread
1	User-level threads are faster to create and manage.	Kernel-level threads are slower to create and manage.
2	Implementation is by a thread library at the user level.	Operating system supports creation of Kernel threads.
3	User-level thread is generic and can run on any operating system.	Kernel-level thread is specific to the operating system.
4	Multi-threaded applications cannot take advantage of multiprocessing.	Kernel routines themselves can be multithreaded.

Implicit Threading

- To transfer the creation and management of threading from application developers to compilers and run-time libraries.
- Thread Pools
- OpenMP
- GCD: Grand Central Dispatch
- ...

Threading Issues

- The `fork()` and `exec()` system calls
 - After calling `fork()`, does the new process duplicates all threads, or is the new process single-threaded?
 - Two variation
 - What about `exec()`?
 - Which variation should be used if `exec()` is called immediately after forking?
 - Duplicating is not necessary!

Thread Cancellation

- Asynchronous cancellation
 - One thread immediately cancels the target thread
 - *Problem*: if resources have been allocated to a cancelled thread or is updating some shared data
- Deferred cancellation
 - The target thread periodically checks whether it should terminate
 - Safe cancellation in *cancellation points*

Signal Handling

- Synchronous signal
 - Delivered to the same process that performed the operation that caused the signal (division by 0)
- Asynchronous signal
 - Generated by an external event (e.g., a timer expiration)
- Signal handlers
 - *Default signal handler* is run by the kernel
 - *User-defined signal handler* overrides default handler

Options for Delivering a Signal

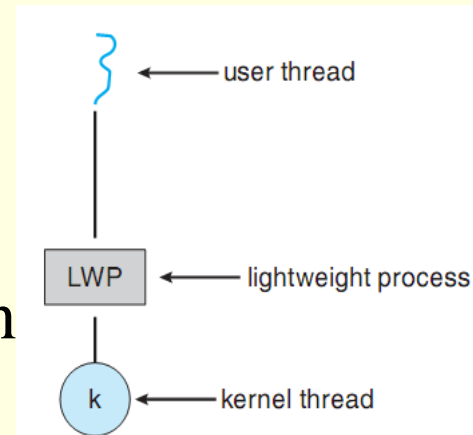
- To the thread to which the signal applies
- To every thread in the process
- To certain threads in the process
- Assign a specific thread to receive all signals to the process

Thread Pools

- Creating a number of threads at process startup and place them into a pool
- Benefits
 - Faster than waiting for creating a thread
 - Limiting the number of threads of each process

Scheduler Activation

- For communication between the kernel and the thread library
 - Many-to-many and two-level models
- LWP: Lightweight Process
 - An intermediate data structure between the user and kernel threads
 - Similar to a virtual processor for the user-thread library
 - Each LWP is attached to a kernel-level thread
 - If the kernel thread blocks, the LWP blocks as well



Number of LWPs

- Assume a CPU-bound application running on a single processor
 - Only one thread can run at once, so one LWP is sufficient
- Assume an IO-bound application
 - One LWP is required for each concurrent blocking system call

How Scheduler Activation Works?

- Kernel informs an application about certain events through *upcalls*
- Upcalls are handled by the thread library with the respective *upcall handler* which is run on a virtual processor
- When a thread is about to block, the kernel makes an upcall and allocates a new virtual processor to the application
- The application runs an upcall handler on this new virtual processor that saves the state of the blocking thread
- The upcall handler schedules another thread on the new virtual processor